

Stop your AI agent from repeating the same mistake

The fix isn't a smarter model. It's a memory file it reads before every task.

THE GAP

No memory = Groundhog Day

You correct it today. Next session it's gone. Most teams never set up CLAUDE.md — so the agent relearns your repo every time.

STEP 01

Know the 4 memory scopes

- Enterprise — org-wide policy
- User — ~/.claude/CLAUDE.md
- Project — ./CLAUDE.md (committed)
- Local — ./CLAUDE.local.md (gitignored)
- Narrowest scope wins on conflict

STEP 02

Bootstrap it in one command

Let Claude read the repo and write a starter file. Then edit it down.

- `/init` → scans repo, drafts CLAUDE.md
- Re-run: it suggests edits, won't overwrite

STEP 03

Keep each file under 200 lines

Anthropic's own rule. Longer files burn context and adherence drops.

- Bad: "Format code properly."
- Good: "Use 2-space indentation."
- Good: "Run npm test before committing."

STEP 04

Add memories without breaking flow

- Start a line with # → saves a memory
- @path/to/file → import another file
- Imports nest up to 4 levels deep

STEP 05

See what's eating your context

- /context → live token breakdown
- /compact → summarize, free up room
- /clear → reset between unrelated tasks

STEP 06

Undo a bad run — code and chat

Every prompt is a checkpoint. Roll back to any of them.

- /rewind (or press Esc twice)
- Restore code, conversation, or both

CHEAT SHEET

Your memory setup in 60 seconds

- /init → draft CLAUDE.md
- Trim to < 200 lines, specific rules
- # ... → add memory mid-task
- /context → audit token usage
- /rewind → undo a bad run



YOUR MOVE

The agent was never the problem

The instructions were. Set the memory once, and it stops guessing. Follow Noblerr for more AI-engineering playbooks.

Follow Noblerr →