

Run OpenAI's terminal coding agent in 60 seconds

Codex reads your repo and edits files from the command line — like Claude Code, different engine.

THE GAP

Most devs let an agent run wild — or not at all

The skill isn't installing it. It's the sandbox + approval policy that lets it work without breaking things.

STEP 01

Install and launch

- `npm install -g @openai/codex`
- `codex`
- Sign in with your ChatGPT account

STEP 02

Set the sandbox dial

How much can it touch on disk?

- read-only → look, don't write
- workspace-write → edit inside project
- danger-full-access → no guardrails

STEP 03

Set the approval dial

When does it stop and ask you?

- untrusted · on-request · never
- Sane default for real work:
- `codex --sandbox workspace-write \`
- `--ask-for-approval on-request`

STEP 04

Go headless for CI & scripts

No interactive prompt. Pipe it anywhere.

- `codex exec "add a CHANGELOG entry`
- `for the latest commit"`
- `Flags: --json --output-last-message`

STEP 05

Wire your tools via MCP

Same servers you already use,
one command.

- `codex mcp add context7 -- \`
- `npx -y @upstash/context7-mcp`
- `codex mcp list` → see what's wired

STEP 06

Give it standing instructions

- AGENTS.md at the repo root
- Read before every task
- Stacks repo-root → subdir (deeper wins)

CHEAT SHEET

Codex from zero to shipping

- `npm install -g @openai/codex`
- `codex --sandbox workspace-write \`
- `--ask-for-approval on-request`
- `codex exec "<task>" → headless`
- `codex mcp add <name> -- <cmd>`



YOUR MOVE

The tool isn't the moat. The workflow is.

We build the AGENTS.md, sandbox policy and MCP wiring so the agent ships instead of guessing. Comment AUTOMATE for the build.

Follow Noblerr →